inovex

inovex classes

Kubernetes

Johannes M. Scheuermann &
Timo Heinrichs

Köln, 27. September 2018

Johannes M. Scheuermann

Macht Dinge in der Wolke

- Team ITO since 2014
- Kubernetes since 2014
- Certified Kubernetes Admin
- LF authorized Kubernetes Admin trainer
- Cloud technologies
- @johscheuer

Timo Heinrichs

Macht Dinge in der Wolke

- Team ITO since 2017
- Docker since v0.9
- Kubernetes since 2017
- Cloud technologies
- @theinrichs

Agenda

- Recap
- Kubernetes concepts
- Volumes and data

- Hands on part: https://github.com/johscheuer/inovex_classes


feedback

What we did

- History of Containers
- History of Kubernetes
- Kubernetes

- Feedback
  - “Micro” Service deployen
  - Autoscaling
  - Fail-Over

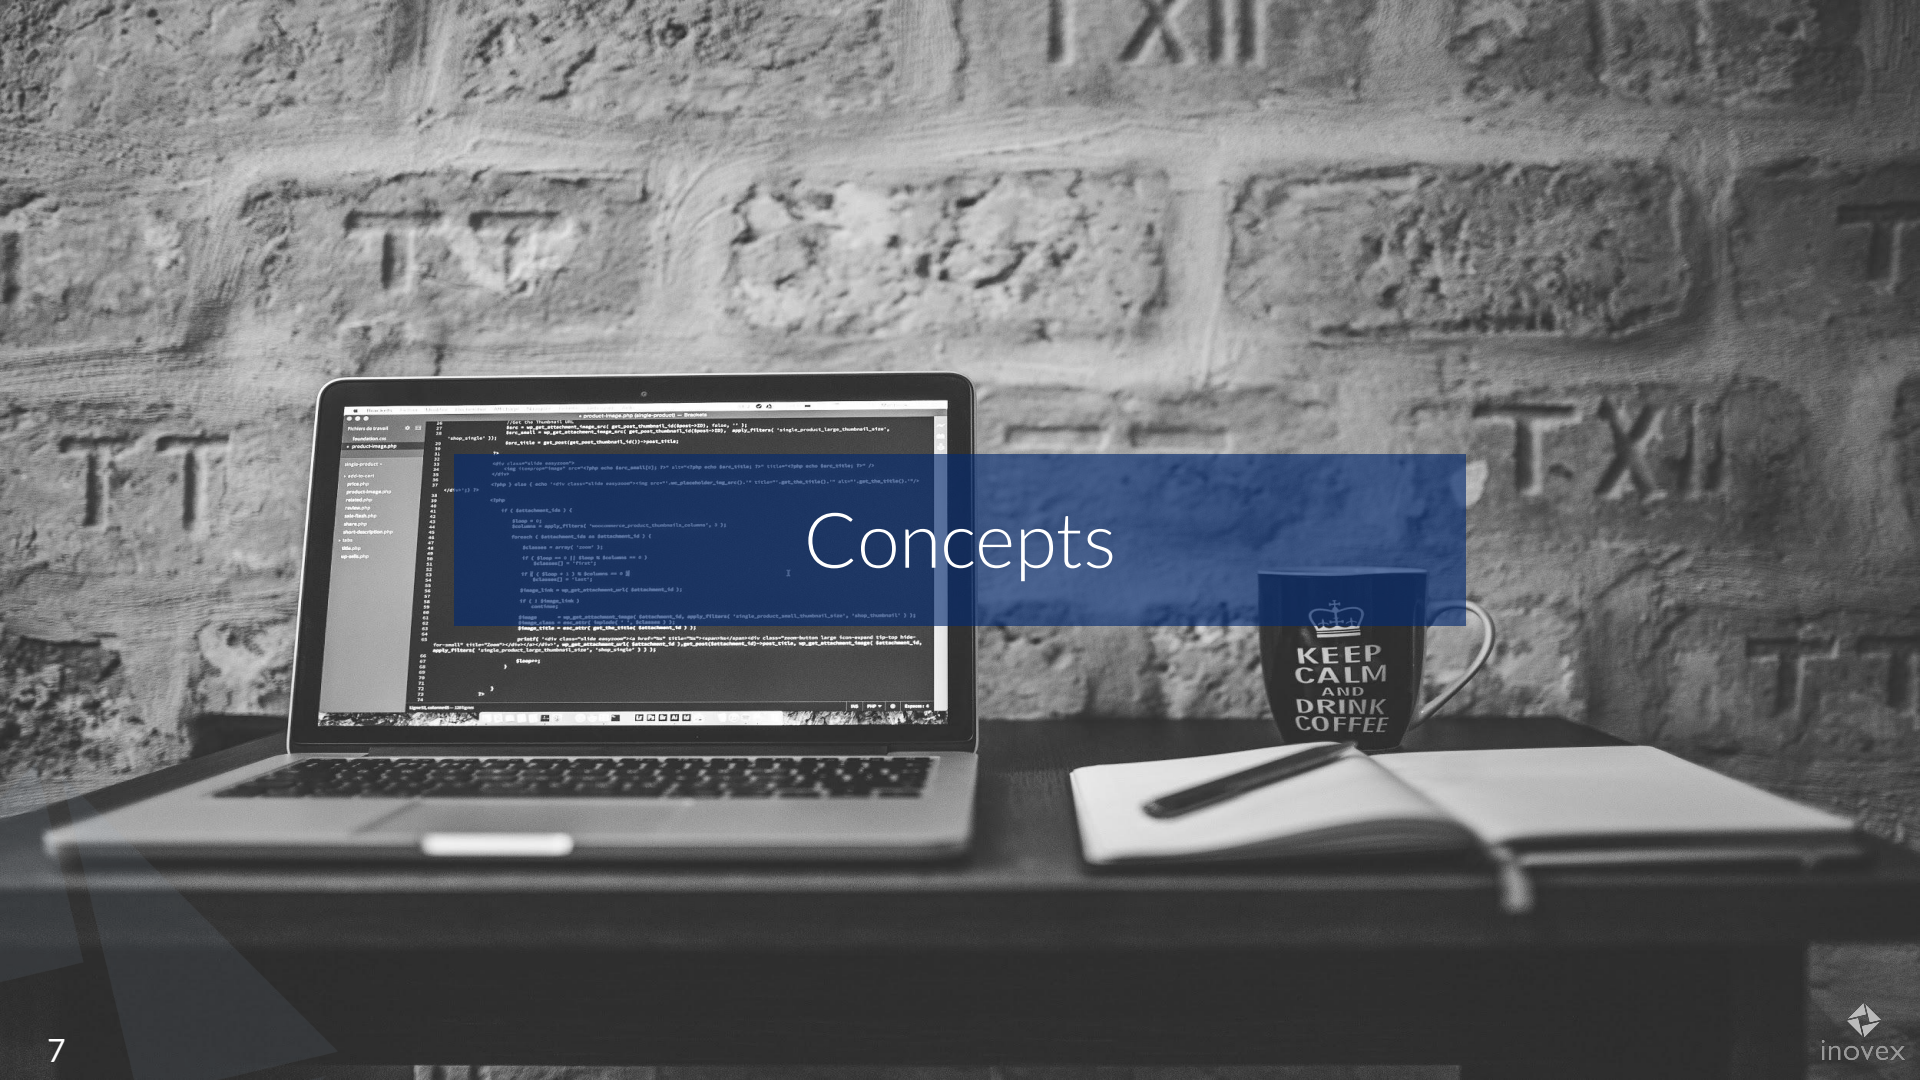

inovex

Kubernetes concepts

- Deployments
- ReplicaSets
- Pods
- Container
- DamonSets
- Jobs
- Services
- Ingress
- ...

Deployments

- High-level abstraction
- Deployments are done on the server-side
- Generates and maintains ReplicaSets
- Works with labels
- Control-loop

ReplicaSet

- Manages one or more Pods
- Control loop
- Ensures that enough replicas are available
- Can be used without Deployments

Pods (recap)

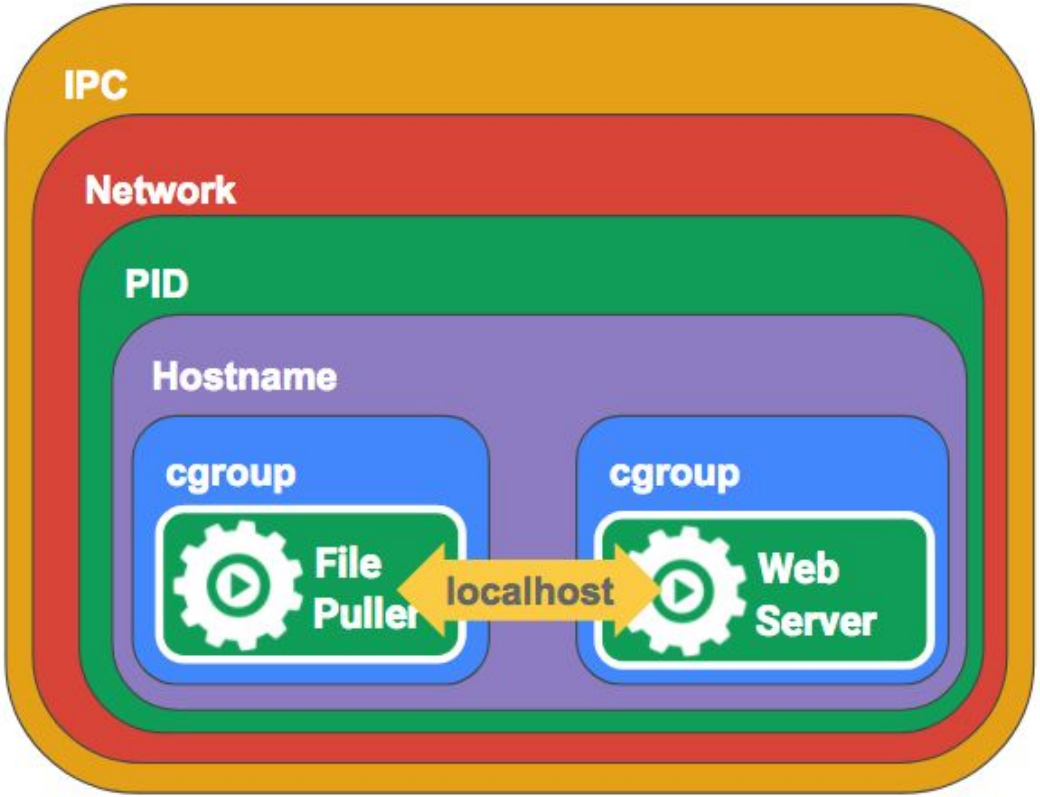

inovex

Rolling Updates

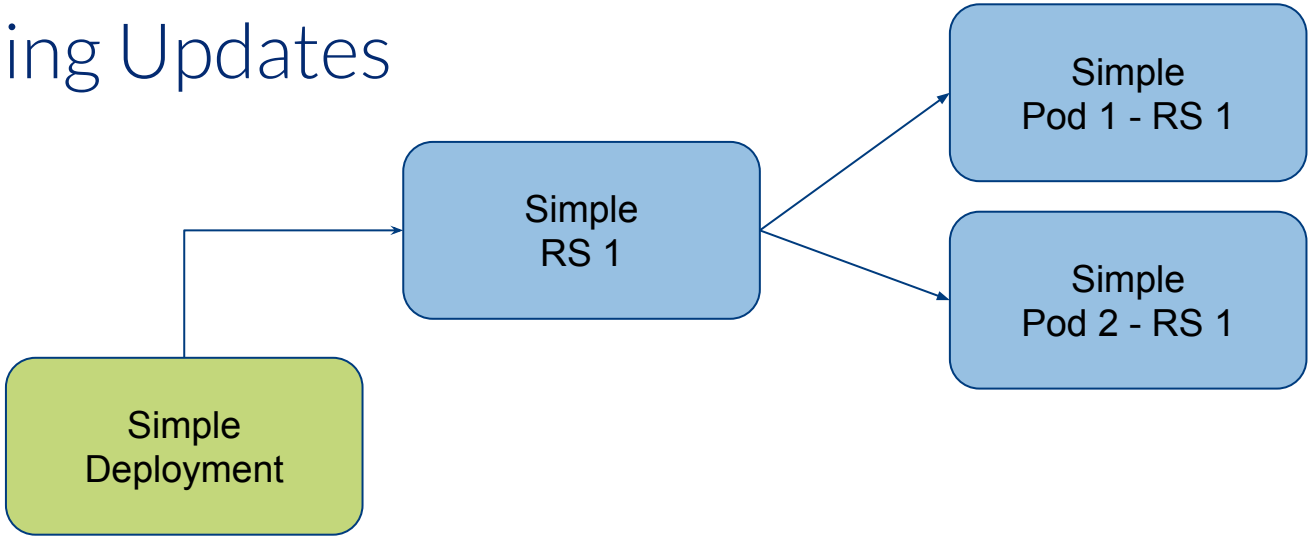

Rolling Updates

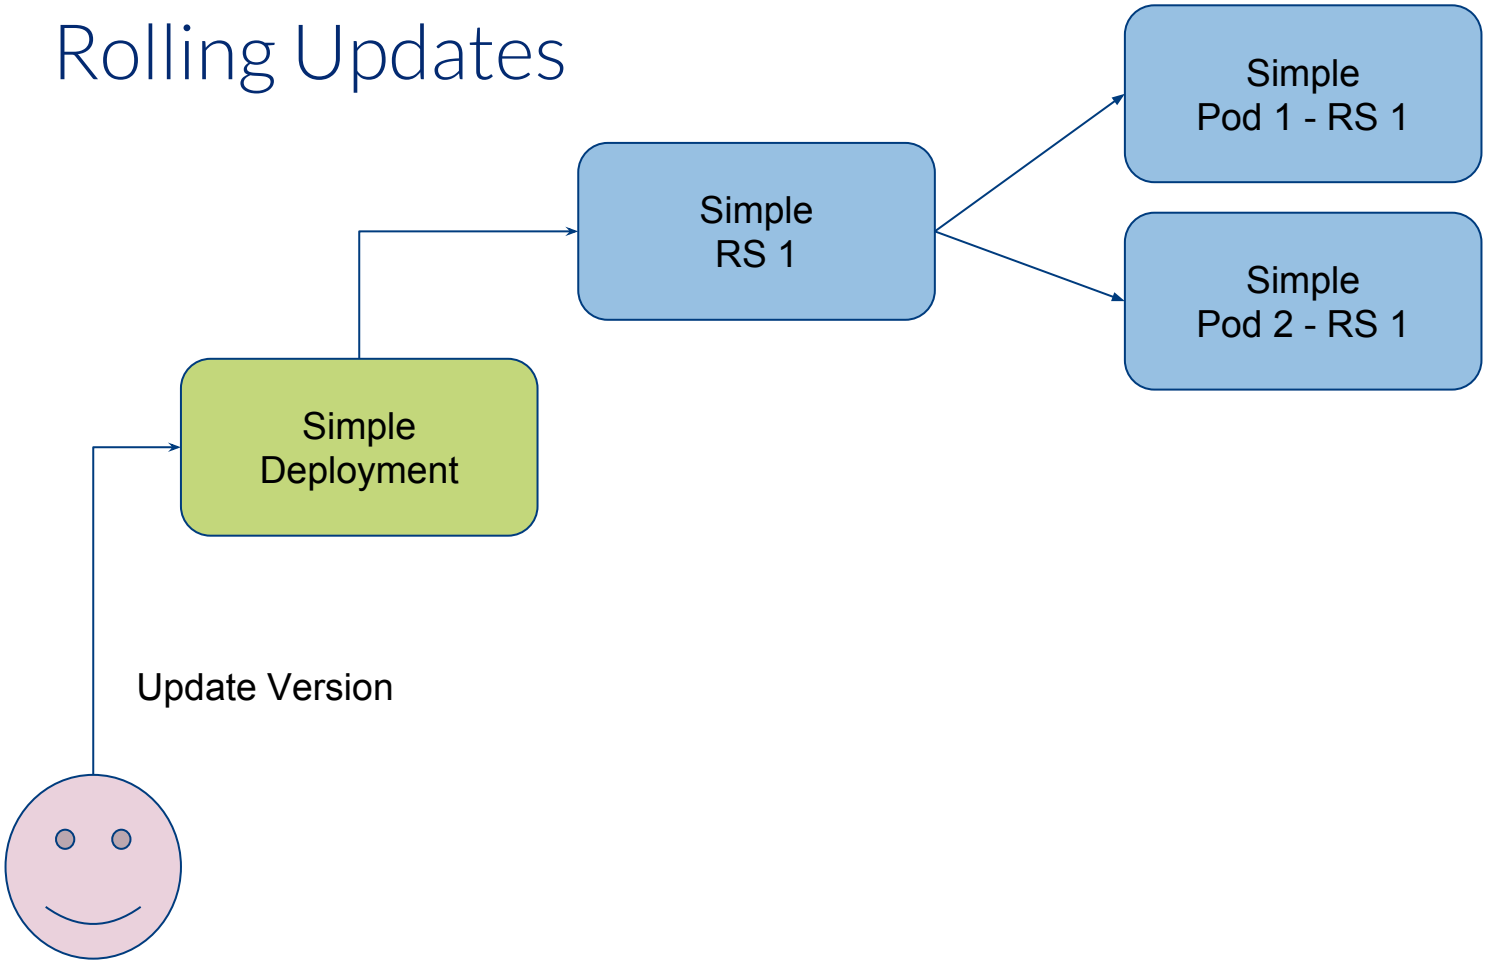

Rolling Updates

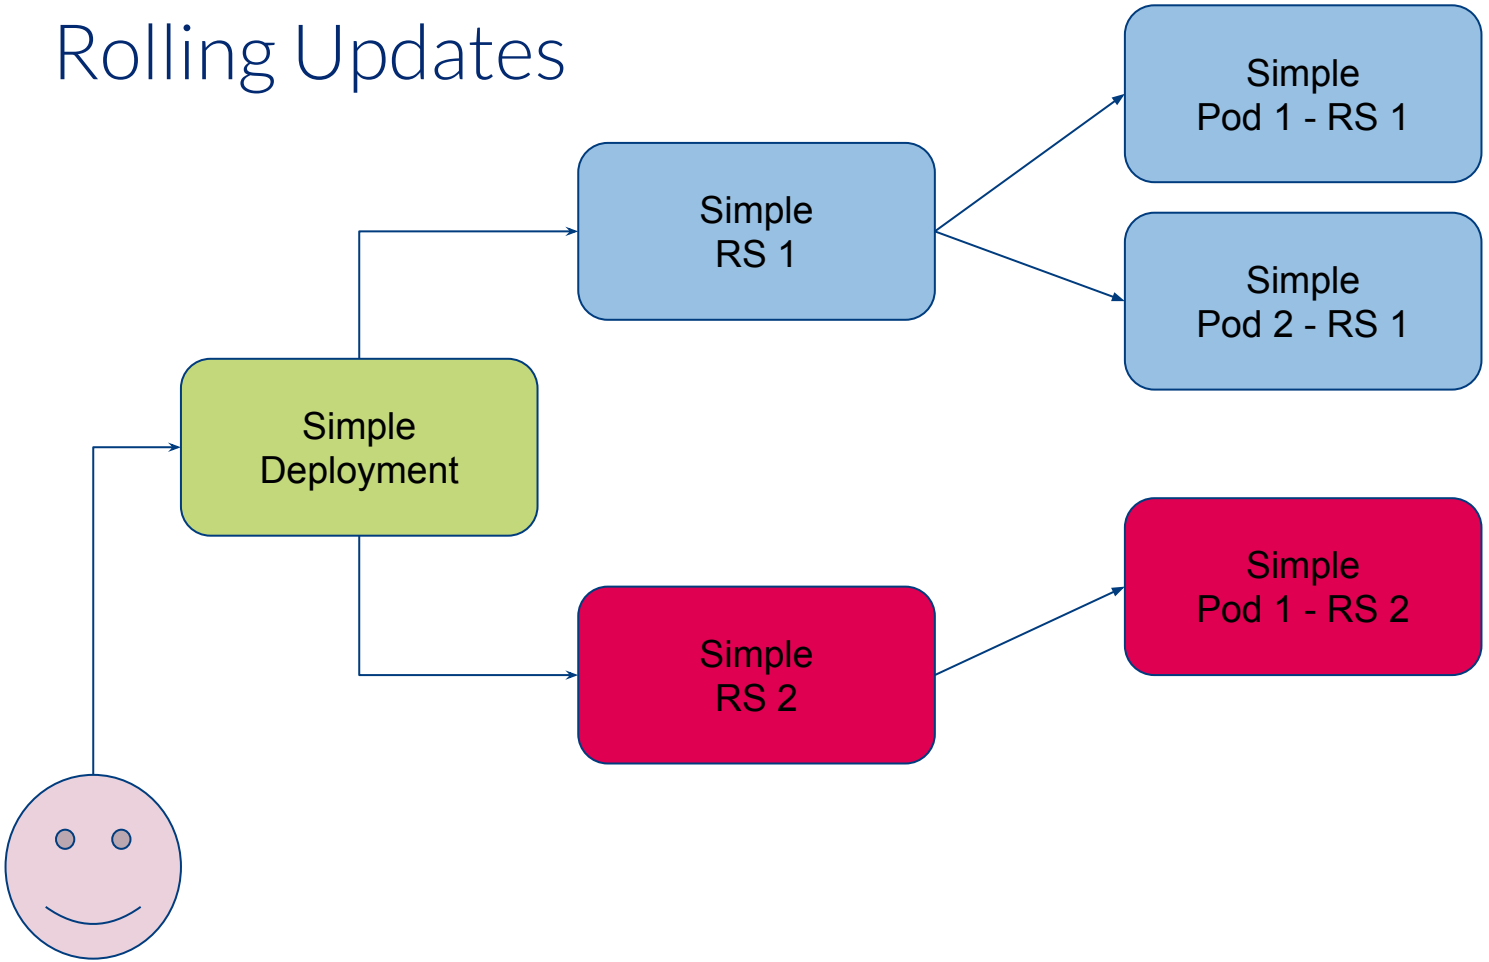

Rolling Updates

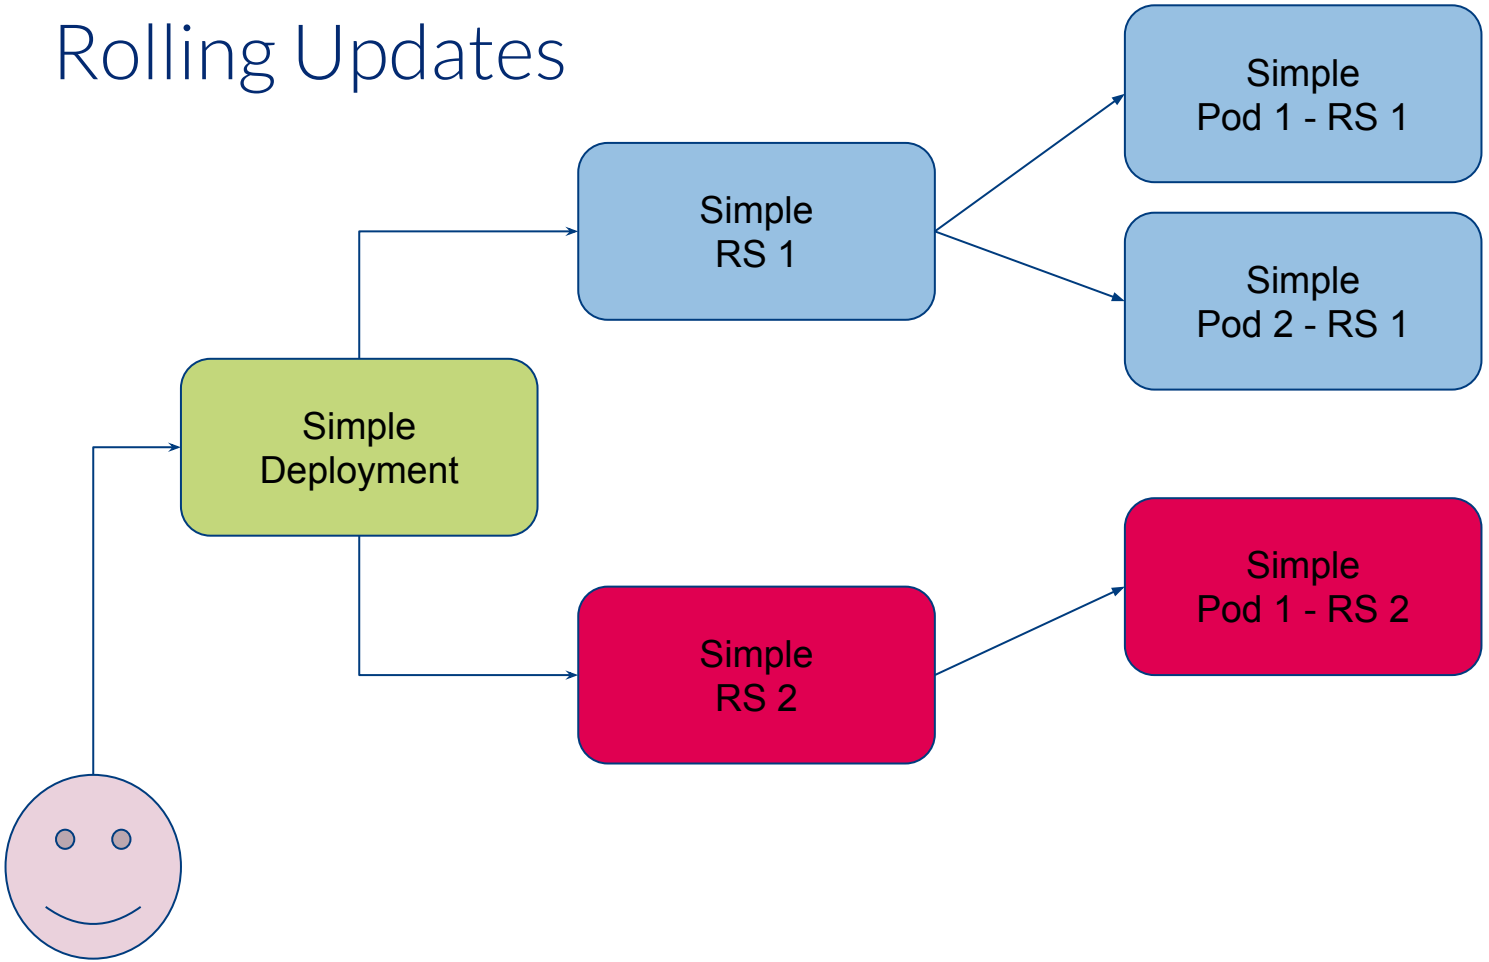

Rolling Updates

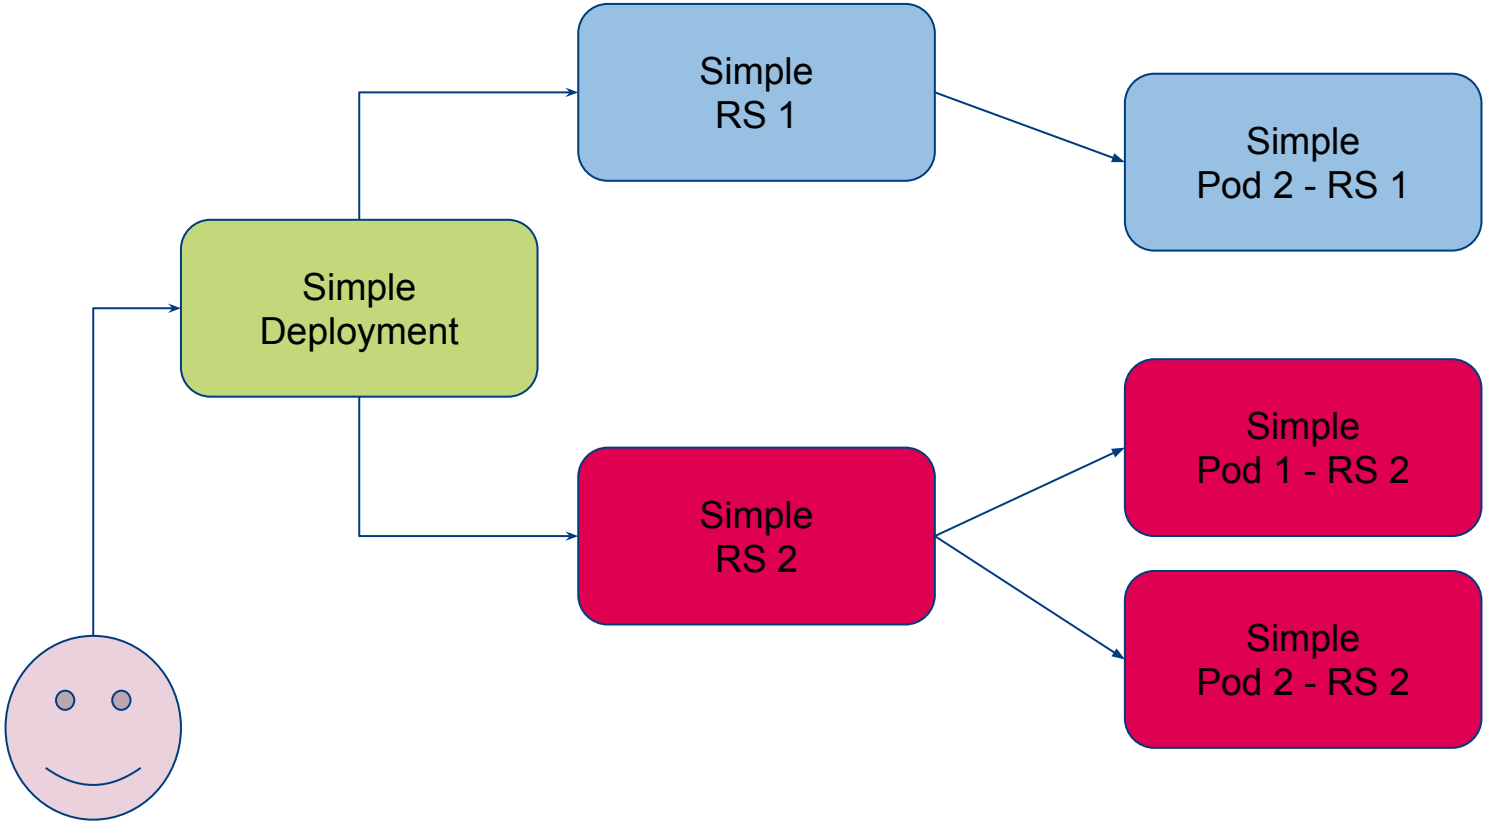

Rolling Updates

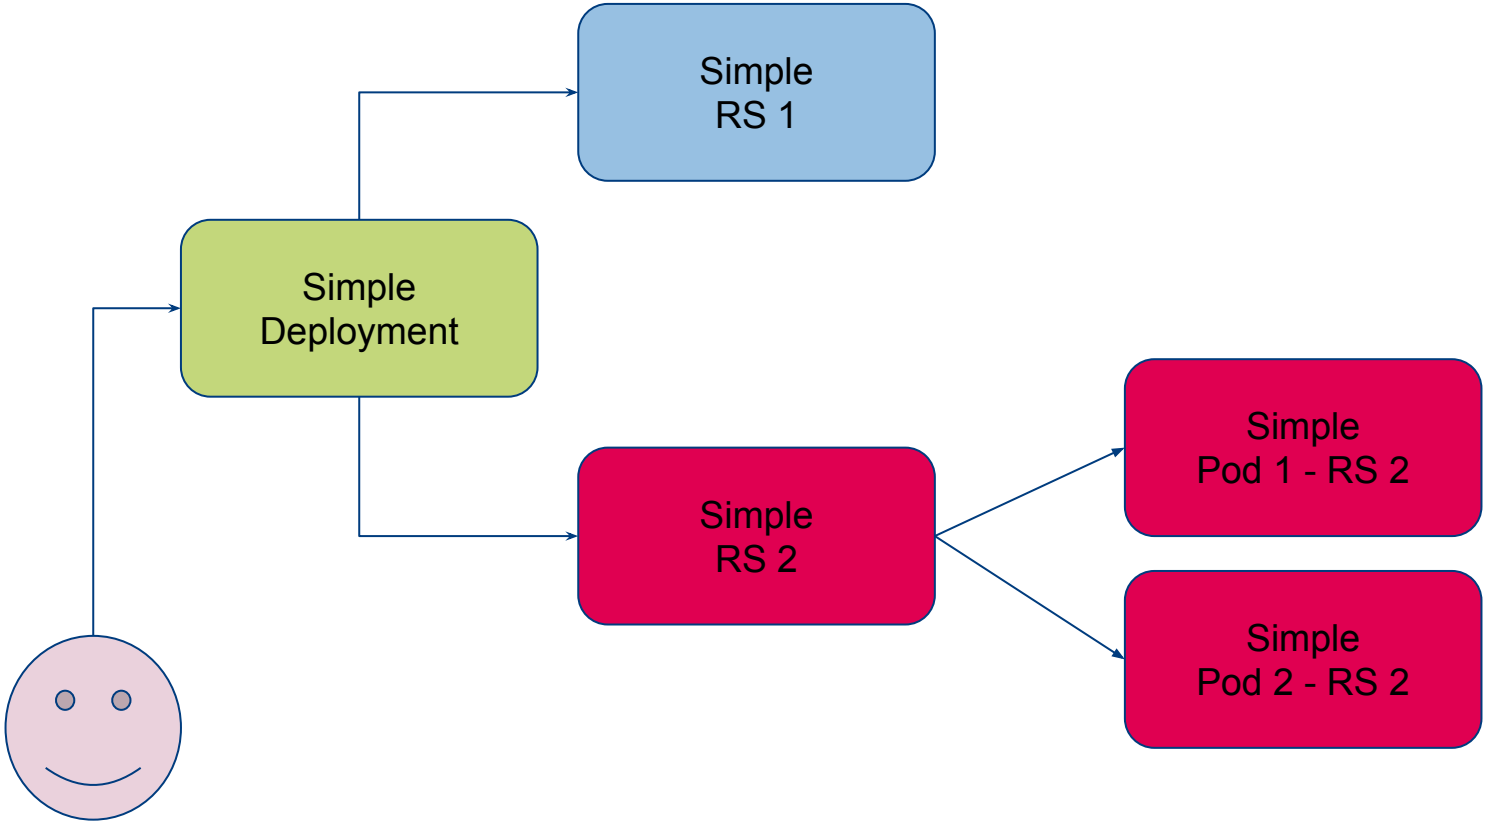

Service

- Provides a stable entrypoint inside the cluster
- Can expose services to the rest of the world
  - Different types
- Select endpoints based on labels
- With a CoreDNS/KubeDNS a DNS entry will be created
  - Only internally

Service

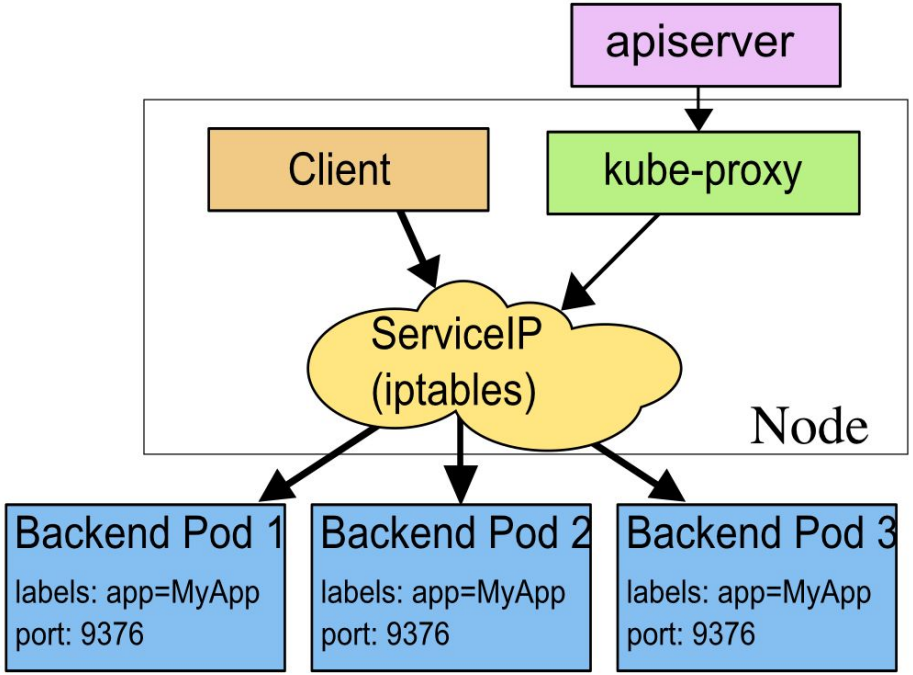

Namespaces

- Logical separation of a cluster
  - e.g. dev/prod
- Can be used for “weak” Multi-Tenancy
- Allows to restrict resource usage (Quotas)
- With RBAC different permissions can be granted

Hands on
Kubernetes concepts

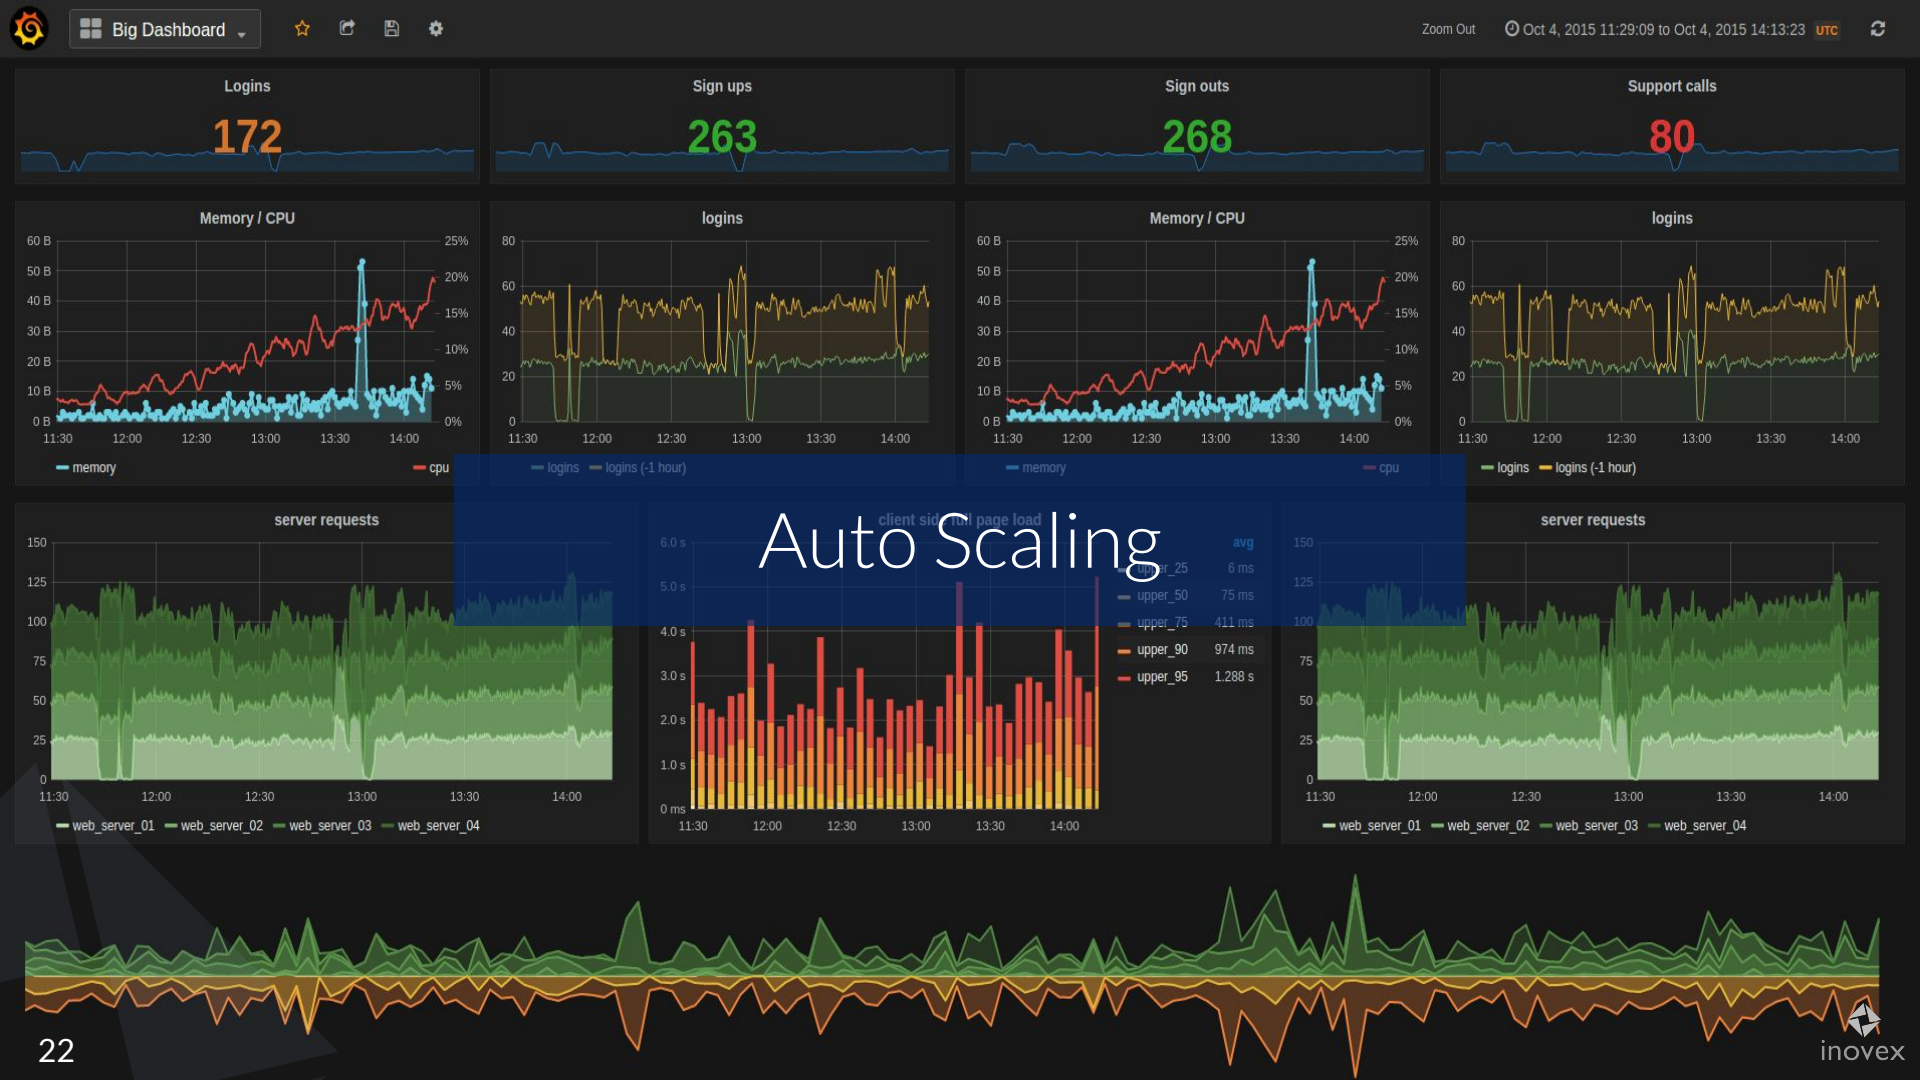

Auto Scaling

inovex

What do we need to auto scale?

Metrics

- Are collected by the “Metrics Server”
- Collects metrics for each node
- Collects metrics for each pod/container
- Needs to be installed additionally
- Works with the API Server Aggregation Layer
  - Metrics Server is an own API server
  - We can access the Metrics Server over the API Server

Auto Scaling

- Deployments (horizontal)
  - Number of Pods based on CPU/Memory
  - Also custom metrics are supported
- Cluster (horizontal)
  - Number of nodes
  - Mostly “dump” algorithms are used
- Pods
  - Vertical Pod Autoscaler (change resource requests/limits)

Horizontal Pod Autoscaler

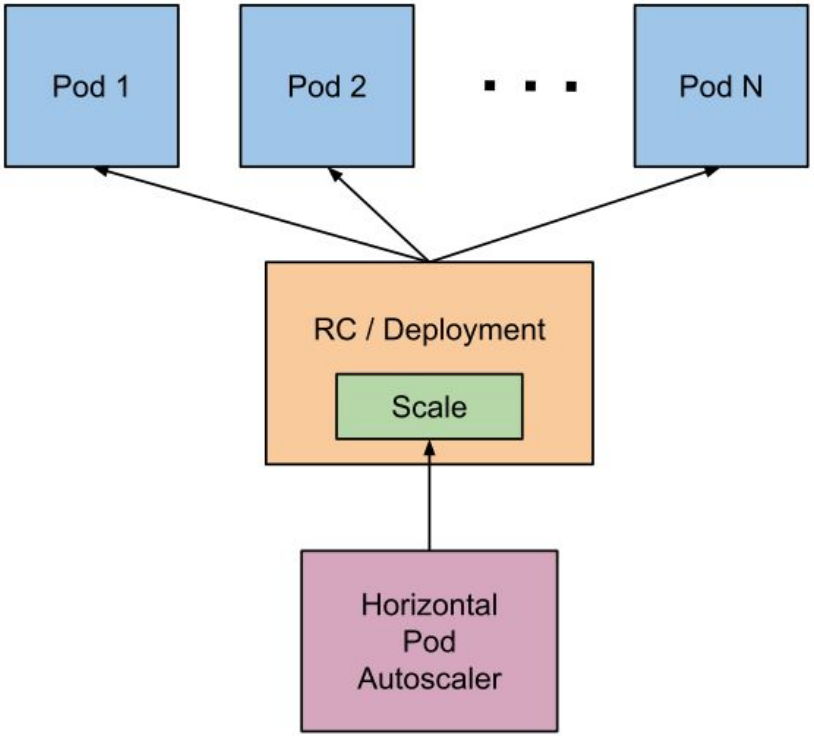

Hands on Autoscaling

Distributed Systems
What could possibly go wrong?

Outlook

- Volumes and Configuration
- Ingress
- Monitoring (Prometheus)

Questions?

Vielen Dank

Johannes M. Scheuermann

inovex GmbH
Ludwig-Erhard-Allee 6
76131 Karlsruhe

jscheuermann@inovex.de
0173 3181 058

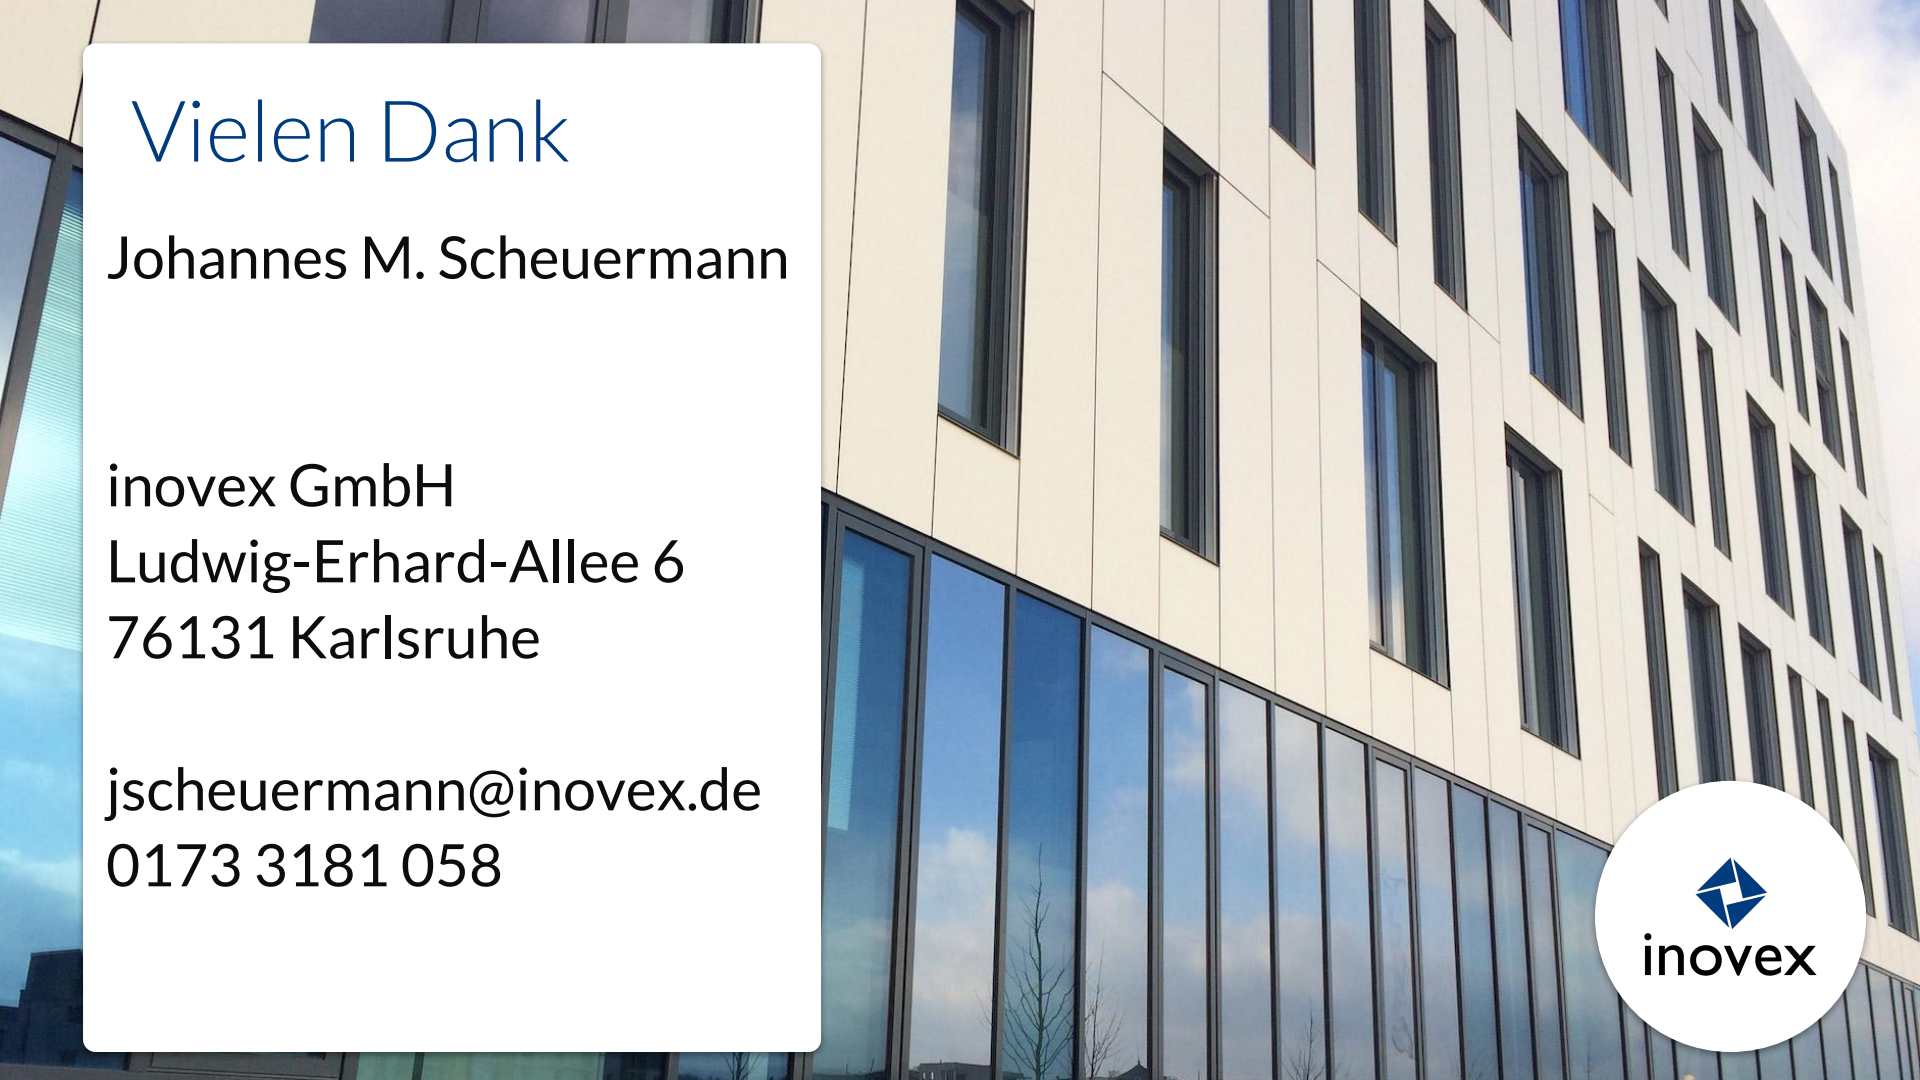

Reading list

- https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale
- https://github.com/kubernetes/community/blob/master/contributors/design-proposals/autoscaling/horizontal-pod-autoscaler.md#autoscaling-algorithm